

# CAS

Compare And Swap (Compare And Exchange) / 自旋 / 自旋锁 / 无锁

因为经常配合循环操作，直到完成为止，所以泛指一类操作

cas(v, a, b)，变量v，期待值a，修改值b

ABA问题，你的女朋友在离开你的这段儿时间经历了别的人，自旋就是你空转等待，一直等到她接纳你为止

解决办法（版本号 AtomicStampedReference），基础类型简单值不需要版本号

## Unsafe

AtomicInteger:

```
1 public final int incrementAndGet() {
2     for (;;) {
3         int current = get();
4         int next = current + 1;
5         if (compareAndSet(current, next))
6             return next;
7     }
8 }
9
10 public final boolean compareAndSet(int expect, int update) {
11     return unsafe.compareAndSwapInt(this, valueOffset, expect, update);
12 }
```

Unsafe:

```
1 public final native boolean compareAndSwapInt(Object var1, long var2, int
  var4, int var5);
```

运用:

```
1 package com.mashibing.jol;
2
3 import sun.misc.Unsafe;
4
5 import java.lang.reflect.Field;
6
7 public class T02_TestUnsafe {
8
9     int i = 0;
10    private static T02_TestUnsafe t = new T02_TestUnsafe();
11
12    public static void main(String[] args) throws Exception {
13        //Unsafe unsafe = Unsafe.getUnsafe();
14
15        Field unsafeField = Unsafe.class.getDeclaredFields()[0];
```

```

16         unsafeField.setAccessible(true);
17         Unsafe unsafe = (Unsafe) unsafeField.get(null);
18
19         Field f = T02_TestUnsafe.class.getDeclaredField("i");
20         long offset = unsafe.objectFieldOffset(f);
21         System.out.println(offset);
22
23         boolean success = unsafe.compareAndSwapInt(t, offset, 0, 1);
24         System.out.println(success);
25         System.out.println(t.i);
26         //unsafe.compareAndSwapInt()
27     }
28 }

```

jdk8u: unsafe.cpp:

cmpxchg = compare and exchange

```

1  UNSAFE_ENTRY(jboolean, Unsafe_CompareAndSwapInt(JNIEnv *env, jobject unsafe,
2  jobject obj, jlong offset, jint e, jint x))
3  Unsafewrapper("Unsafe_CompareAndSwapInt");
4  oop p = JNIHandles::resolve(obj);
5  jint* addr = (jint *) index_oop_from_field_offset_long(p, offset);
6  return (jint)(Atomic::cmpxchg(x, addr, e)) == e;
7  UNSAFE_END

```

jdk8u: atomic\_linux\_x86.inline.hpp

is\_MP = Multi Processor

```

1  inline jint    Atomic::cmpxchg    (jint    exchange_value, volatile jint*
2  dest, jint    compare_value) {
3  int mp = os::is_MP();
4  __asm__ volatile (LOCK_IF_MP(%4) "cmpxchgl %1,(%3)"
5  : "=a" (exchange_value)
6  : "r" (exchange_value), "a" (compare_value), "r" (dest),
7  "r" (mp)
8  : "cc", "memory");
9  return exchange_value;
10 }

```

jdk8u: os.hpp is\_MP()

```

1  static inline bool is_MP() {
2  // During bootstrap if _processor_count is not yet initialized
3  // we claim to be MP as that is safest. If any platform has a
4  // stub generator that might be triggered in this phase and for
5  // which being declared MP when in fact not, is a problem - then
6  // the bootstrap routine for the stub generator needs to check
7  // the processor count directly and leave the bootstrap routine
8  // in place until called after initialization has occurred.
9  return (_processor_count != 1) || AssumeMP;
10 }

```

jdk8u: atomic\_linux\_x86.inline.hpp

```
1 | #define LOCK_IF_MP(mp) "cmp $0, " #mp "; je 1f; lock; 1: "
```

最终实现:

cmpxchg = cas修改变量值

```
1 | lock cmpxchg 指令
```

硬件:

lock指令在执行后面指令的时候锁定一个北桥信号

(不采用锁总线的方式)

## markword

## 工具: JOL = Java Object Layout

```
1 | <dependencies>
2 |     <!-- https://mvnrepository.com/artifact/org.openjdk.jol/jol-core -->
3 |     <dependency>
4 |         <groupId>org.openjdk.jol</groupId>
5 |         <artifactId>jol-core</artifactId>
6 |         <version>0.9</version>
7 |     </dependency>
8 | </dependencies>
```

jdk8u: markOop.hpp

```
1 | // Bit-format of an object header (most significant first, big endian
  | // layout below):
2 | //
3 | // 32 bits:
4 | // -----
5 | //          hash:25 ----->| age:4    biased_lock:1 lock:2 (normal
  | // object)
6 | //          JavaThread*:23 epoch:2 age:4    biased_lock:1 lock:2 (biased
  | // object)
7 | //          size:32 ----->| (CMS
  | // free block)
8 | //          PromotedObject*:29 ----->| promo_bits:3 ----->| (CMS
  | // promoted object)
9 | //
10 | // 64 bits:
11 | // -----
12 | // unused:25 hash:31 -->| unused:1    age:4    biased_lock:1 lock:2 (normal
  | // object)
13 | // JavaThread*:54 epoch:2 unused:1    age:4    biased_lock:1 lock:2 (biased
  | // object)
```

```

14 // PromotedObject*:61 ----->| promo_bits:3 ----->| (CMS
    promoted object)
15 // size:64 ----->| (CMS
    free block)
16 //
17 // unused:25 hash:31 -->| cms_free:1 age:4    biased_lock:1 lock:2 (COOPs
    && normal object)
18 // JavaThread*:54 epoch:2 cms_free:1 age:4    biased_lock:1 lock:2 (COOPs
    && biased object)
19 // narrowOop:32 unused:24 cms_free:1 unused:4 promo_bits:3 ----->| (COOPs
    && CMS promoted object)
20 // unused:21 size:35 -->| cms_free:1 unused:7 ----->| (COOPs
    && CMS free block)

```

## synchronized的横切面详解

---

1. synchronized原理
2. 升级过程
3. 汇编实现
4. vs reentrantLock的区别

## java源码层级

---

synchronized(o)

## 字节码层级

---

monitorenter moniterexit

## JVM层级 (Hotspot)

---

```

1 package com.mashibing.insidesync;
2
3 import org.openjdk.jol.info.ClassLayout;
4
5 public class T01_Sync1 {
6
7
8     public static void main(String[] args) {
9         Object o = new Object();
10
11         System.out.println(ClassLayout.parseInstance(o).toPrintable());
12     }
13 }

```

```

1 com.mashibing.insidesync.T01_Sync1$Lock object internals:
2   OFFSET  SIZE   TYPE DESCRIPTION                               VALUE
3     0      4   (object header) 05 00 00 00 (00000101 00000000 00000000
00000000) (5)
4     4      4   (object header) 00 00 00 00 (00000000 00000000 00000000
00000000) (0)
5     8      4   (object header) 49 ce 00 20 (01001001 11001110 00000000
00100000) (536923721)
6    12      4               (loss due to the next object alignment)
7 Instance size: 16 bytes
8 Space losses: 0 bytes internal + 4 bytes external = 4 bytes total

```

```

1 com.mashibing.insidesync.T02_Sync2$Lock object internals:
2   OFFSET  SIZE   TYPE DESCRIPTION                               VALUE
3     0      4   (object header) 05 90 2e 1e (00000101 10010000 00101110
00011110) (506368005)
4     4      4   (object header) 1b 02 00 00 (00011011 00000010 00000000
00000000) (539)
5     8      4   (object header) 49 ce 00 20 (01001001 11001110 00000000
00100000) (536923721)
6    12      4               (loss due to the next object alignment)
7 Instance size: 16 bytes
8 Space losses: 0 bytes internal + 4 bytes external = 4 bytes total

```

InterpreterRuntime::monitorenter方法

```

1 IRT_ENTRY_NO_ASYNC(void, InterpreterRuntime::monitorenter(JavaThread*
thread, BasicObjectLock* elem))
2 #ifdef ASSERT
3   thread->last_frame().interpreter_frame_verify_monitor(elem);
4 #endif
5   if (PrintBiasedLockingStatistics) {
6     Atomic::inc(BiasedLocking::slow_path_entry_count_addr());
7   }
8   Handle h_obj(thread, elem->obj());
9   assert(Universe::heap()->is_in_reserved_or_null(h_obj()),
10          "must be NULL or an object");
11   if (UseBiasedLocking) {
12     // Retry fast entry if bias is revoked to avoid unnecessary inflation
13     ObjectsSynchronizer::fast_enter(h_obj, elem->lock(), true, CHECK);
14   } else {
15     ObjectsSynchronizer::slow_enter(h_obj, elem->lock(), CHECK);
16   }
17   assert(Universe::heap()->is_in_reserved_or_null(elem->obj()),
18          "must be NULL or an object");
19   #ifdef ASSERT
20     thread->last_frame().interpreter_frame_verify_monitor(elem);
21   #endif
22   IRT_END

```

synchronizer.cpp

revoke\_and\_rebias

```

1 void ObjectSynchronizer::fast_enter(Handle obj, BasicLock* lock, bool
  attempt_rebias, TRAPS) {
2   if (UseBiasedLocking) {
3     if (!SafepointSynchronize::is_at_safepoint()) {
4       BiasedLocking::Condition cond = BiasedLocking::revoke_and_rebias(obj,
  attempt_rebias, THREAD);
5       if (cond == BiasedLocking::BIAS_REVOKED_AND_REBIASED) {
6         return;
7       }
8     } else {
9       assert(!attempt_rebias, "can not rebias toward VM thread");
10      BiasedLocking::revoke_at_safepoint(obj);
11    }
12    assert(!obj->mark()->has_bias_pattern(), "biases should be revoked by
  now");
13  }
14
15  slow_enter (obj, lock, THREAD) ;
16 }

```

```

1 void ObjectSynchronizer::slow_enter(Handle obj, BasicLock* lock, TRAPS) {
2   markOop mark = obj->mark();
3   assert(!mark->has_bias_pattern(), "should not see bias pattern here");
4
5   if (mark->is_neutral()) {
6     // Anticipate successful CAS -- the ST of the displaced mark must
7     // be visible <= the ST performed by the CAS.
8     lock->set_displaced_header(mark);
9     if (mark == (markOop) Atomic::cmpxchg_ptr(lock, obj()->mark_addr(),
  mark)) {
10      TEVENT (slow_enter: release stacklock) ;
11      return ;
12    }
13    // Fall through to inflate() ...
14  } else
15    if (mark->has_locker() && THREAD->is_lock_owned((address)mark->locker()))
  {
16      assert(lock != mark->locker(), "must not re-lock the same lock");
17      assert(lock != (BasicLock*)obj->mark(), "don't relock with same
  BasicLock");
18      lock->set_displaced_header(NULL);
19      return;
20    }
21
22  #if 0
23    // The following optimization isn't particularly useful.
24    if (mark->has_monitor() && mark->monitor()->is_entered(THREAD)) {
25      lock->set_displaced_header (NULL) ;
26      return ;
27    }
28  #endif
29
30  // The object header will never be displaced to this lock,
31  // so it does not matter what the value is, except that it
32  // must be non-zero to avoid looking like a re-entrant lock,
33  // and must not look locked either.
34  lock->set_displaced_header(markOopDesc::unused_mark());

```

```
35 | ObjectSynchronizer::inflate(THREAD, obj())->enter(THREAD);
36 | }
```

inflate方法：膨胀为重量级锁

## 锁升级过程

### JDK8 markword实现表：



无锁 - 偏向锁 - 轻量级锁（自旋锁，自适应自旋） - 重量级锁

synchronized优化的过程和markword息息相关

用markword中最低的三位代表锁状态 其中1位是偏向锁位 两位是普通锁位

1. Object o = new Object()

锁 = 0 01 无锁态

2. o.hashCode()

001 + hashCode

```
1 | 00000001 10101101 00110100 00110110
2 | 01011001 00000000 00000000 00000000
```

little endian big endian

00000000 00000000 00000000 01011001 00110110 00110100 10101101 00000000

3. 默认synchronized(o)

00 -> 轻量级锁

默认情况 偏向锁有个时延，默认是4秒

why? 因为JVM虚拟机自己有一些默认启动的线程，里面有好多sync代码，这些sync代码启动时就知道肯定会有竞争，如果使用偏向锁，就会造成偏向锁不断的进行锁撤销和锁升级的操作，效率较低。

```
1 | -XX:BiasedLockingStartupDelay=0
```

4. 如果设定上述参数

new Object () -> 101 偏向锁 -> 线程ID为0 -> Anonymous BiasedLock

打开偏向锁，new出来的对象，默认就是一个可偏向匿名对象101

5. 如果有线程上锁

上偏向锁，指的就是，把markword的线程ID改为自己线程ID的过程

偏向锁不可重偏向 批量偏向 批量撤销

6. 如果有线程竞争

撤销偏向锁，升级轻量级锁

线程在自己的线程栈生成LockRecord，用CAS操作将markword设置为指向自己这个线程的LR的指针，设置成功者得到锁

## 7. 如果竞争加剧

竞争加剧：有线程超过10次自旋， `-XX:PreBlockSpin`， 或者自旋线程数超过CPU核数的一半，  
1.6之后，加入自适应自旋 `Adaptive Self Spinning`， JVM自己控制

升级重量级锁：-> 向操作系统申请资源， linux mutex，CPU从3级-0级系统调用，线程挂起，进入等待队列，等待操作系统的调度，然后再映射回用户空间

(以上实验环境是JDK11，打开就是偏向锁，而JDK8默认对象头是无锁)

偏向锁默认是打开的，但是有一个时延，如果要观察到偏向锁，应该设定参数

没错，我就是厕所所长

加锁，指的是锁定对象

锁升级的过程

JDK较早的版本 OS的资源 互斥量 用户态 -> 内核态的转换 重量级 效率比较低

现代版本进行了优化

无锁 - 偏向锁 -轻量级锁（自旋锁）-重量级锁

偏向锁 - markword 上记录当前线程指针，下次同一个线程加锁的时候，不需要争用，只需要判断线程指针是否同一个，所以，偏向锁，偏向加锁的第一个线程。hashCode备份在线程栈上 线程销毁，锁降级为无锁

有争用 - 锁升级为轻量级锁 - 每个线程有自己的LockRecord在自己的线程栈上，用CAS去争用markword的LR的指针，指针指向哪个线程的LR，哪个线程就拥有锁

自旋超过10次，升级为重量级锁 - 如果太多线程自旋 CPU消耗过大，不如升级为重量级锁，进入等待队列（不消耗CPU） `-XX:PreBlockSpin`

自旋锁在JDK1.4.2中引入，使用 `-XX:+UseSpinning` 来开启。JDK 6 中变为默认开启，并且引入了自适应的自旋锁（适应性自旋锁）。

自适应自旋锁意味着自旋的时间（次数）不再固定，而是由前一次在同一个锁上的自旋时间及锁的拥有者的状态来决定。如果在同一个锁对象上，自旋等待刚刚成功获得过锁，并且持有锁的线程正在运行中，那么虚拟机就会认为这次自旋也是很有可能再次成功，进而它将允许自旋等待持续相对更长的时间。如果对于某个锁，自旋很少成功获得过，那在以后尝试获取这个锁时将可能省略掉自旋过程，直接阻塞线程，避免浪费处理器资源。

偏向锁由于有锁撤销的过程revoke，会消耗系统资源，所以，在锁争用特别激烈的时候，用偏向锁未必效率高。还不如直接使用轻量级锁。

## synchronized最底层实现

```
1
2 public class T {
3     static volatile int i = 0;
4
5     public static void n() { i++; }
6
7     public static synchronized void m() {}
```



```

8
9     public static void main(String[] args) {
10         for(int j=0; j<1000_000; j++) {
11             m();
12             n();
13         }
14     }
15 }
16

```

java -XX:+UnlockDiagnosticVMOptions -XX:+PrintAssembly T

C1 Compile Level 1 (一级优化)

C2 Compile Level 2 (二级优化)

找到m() n()方法的汇编码，会看到 lock comxchg .....指令

## synchronized vs Lock (CAS)

- 1 在高争用 高耗时的环境下synchronized效率更高
- 2 在低争用 低耗时的环境下CAS效率更高
- 3 synchronized到重量级之后是等待队列（不消耗CPU）
- 4 CAS（等待期间消耗CPU）
- 5
- 6 一切以实测为准

## 锁消除 lock eliminate

```

1 public void add(String str1,String str2){
2     StringBuffer sb = new StringBuffer();
3     sb.append(str1).append(str2);
4 }

```

我们都知道 StringBuffer 是线程安全的，因为它的关键方法都是被 synchronized 修饰过的，但我们看上面这段代码，我们会发现，sb 这个引用只会在 add 方法中使用，不可能被其它线程引用（因为是局部变量，栈私有），因此 sb 是不可能共享的资源，JVM 会自动消除 StringBuffer 对象内部的锁。

## 锁粗化 lock coarsening

```

1 public String test(String str){
2
3     int i = 0;
4     StringBuffer sb = new StringBuffer();
5     while(i < 100){
6         sb.append(str);
7         i++;
8     }
9     return sb.toString();
10 }

```

JVM 会检测到这样一连串的操作都对同一个对象加锁（while 循环内 100 次执行 append，没有锁粗化的就要进行 100 次加锁/解锁），此时 JVM 就会将加锁的范围粗化到这一连串的操作的外部（比如 while 虚幻体外），使得这一连串操作只需要加一次锁即可。

## 锁降级（不重要）

---

<https://www.zhihu.com/question/63859501>

其实，只被VMThread访问，降级也就没啥意义了。所以可以简单认为锁降级不存在！

## 超线程

---

一个ALU + 两组Registers + PC

## 参考资料

---

<http://openjdk.java.net/groups/hotspot/docs/HotSpotGlossary.html>